



diagonals: Fat Diagonals in R

Bastiaan Quast

The Graduate Institute, Geneva

Abstract

The abstract of the article.

Keywords: matrix, diagonal, block, R.

1. Introduction

Diagonals are an important matrix manipulation. We present the `diagonals` R package, which implements functions for dealing with **fat diagonals**. Fat diagonals are block matrix-diagonals that occur when two or more dimensions are mapped along a single edge of a matrix. For an asymmetric network graph (e.g. a dyadic social network) to be mapped to a matrix, we would need each node along each edge of matrix, however we would also need to map the direction of the tie, which is an additional dimension. Typically these would be represented as higher-order arrays (i.e. ≥ 3). In order to effectively visualise such arrays, it can be helpful to do so in a matrix. This could for instance be represented as in the following matrix (where the `i` and `o` suffices represent incoming and outgoing receptively).

	Ai	Ao	Bi	Bo	Ci	Co	Di	Do
A	1	1	0	0	1	0	1	0
B	0	1	1	1	0	1	0	1
C	1	0	0	0	1	1	0	0
D	1	0	1	0	1	1	1	1

Sometimes the ties of a node to itself are not particularly meaningful (e.g. feeling of amiability towards oneself) and can be removed. For a symmetric network this can simply be done using the function `diag()` in R's `base` package, e.g.

```
sm <- matrix(1, nrow=4, ncol=4)
```

```
diag(sm) <- NA
sm
```

```
      [,1] [,2] [,3] [,4]
[1,]    NA     1     1     1
[2,]     1    NA     1     1
[3,]     1     1    NA     1
[4,]     1     1     1    NA
```

However, for higher-order matrices this does not work well.

```
diag(m) <- NA
m
```

```
      Ai Ao Bi Bo Ci Co Di Do
A NA  1  0  0  1  0  1  0
B  0 NA  1  1  0  1  0  1
C  1  0 NA  0  1  1  0  0
D  1  0  1 NA  1  1  1  1
```

In comes the **diagonals** package and its workhorse **fatdiag()** function. The function is designed to mimic the behaviour of the **diag()** as closely as possible, but with then for **fat diagonals**.

```
library(diagonals)
```

```
# the matrix m was restored to its original state
fatdiag(m, steps=4) <- NA
m
```

```
      Ai Ao Bi Bo Ci Co Di Do
A NA NA  0  0  1  0  1  0
B  0  1 NA NA  0  1  0  1
C  1  0  0  0 NA NA  0  0
D  1  0  1  0  1  1 NA NA
```

Note that the **steps** argument defines the number of steps on the diagonal ladder. Alternatively we could set the **size** of the step, more on this later.

2. Design

The implementation of fat diagonals in the **diagonals** package is intended to be as close as possible to the functions dealing with diagonals included in the **base** package. As such, the package includes two functions.

- **fatdiag()**

- `fatdiag()`<-

These functions offer a very similar syntax to the base functions:

- `diag()`
- `diag()`<-

With the exception that the fat diagonal functions generally need more information in terms of the number of **steps** on the diagonal ladder, or the **size** of these steps.

The function `fatdiag`<- like its base package equivalent replaces the (fat) diagonal of its first argument `x` with the right side argument `value`. The `value` argument can either be a scalar, in which case it is recycled for the length of the diagonal, or it can be vector. For the base package function `diag()`<- this vector has be of the same length as the diagonal (here, the shortest dimension of the matrix), however, the `fatdiag()`<- function will accept any vector that is of a length that is an integer divisor of the length of the diagonal. For example, if the length of the diagonal is 12, then the follow lengths for the replacement vector are accepted: 1, 2, 3, 4, and 6.

The `fatdiag` function act similar to the `diag()` function. Both these functions have two main applications . The first application is (fat) diagonal extraction, is the first argument `x` is a matrix, i.e. `length(dim(x)) == 2`, then the function extracts the diagonal matrix and returns it as a vector.

The second application is (fat) diagonal matrix creation. This can be done in two ways, using a scalar, or using a vector. If a scalar is used for `x`, the `diag()` function returns an identity matrix `Ix`, i.e. a matrix of dimensions `x` times `x` is returned, with 1 on the diagonal positions and 0 elsewhere. The `fatdiag()` function supports the creation of non-square matrices (e.g. using `size = c(3,2)`) and therefore uses `x` as the longest dimension of the matrix, where the other dimension is determined automatically using the `size` argument.

3. Usage

In the introduction we briefly demonstrate the usage of the `fatdiag()` function for assigning a new `value` to the fat diagonal. Here we take a closer look at some of the additional functionality that is implemented.

```
fatdiag(m, size=c(1,2) ) <- 881:888
m
```

	Ai	Ao	Bi	Bo	Ci	Co	Di	Do
A	881	882	0	0	1	0	1	0
B	0	1	883	884	0	1	0	1
C	1	0	0	0	885	886	0	0
D	1	0	1	0	1	1	887	888

So far we have been using the set `fatdiag()`, i.e. `fatdiag()`<-. However, we can also use the `fatdiag()` function either for diagonal extraction, or diagonal matrix creation.

```
fatdiag(m, steps = 4)
```

```
[1] 881 882 883 884 885 886 887 888
```

Fat diagonal matrices can be created using a scalar:

```
fatdiag(9, steps=3)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
[1,]	1	1	1	0	0	0	0	0	0
[2,]	1	1	1	0	0	0	0	0	0
[3,]	1	1	1	0	0	0	0	0	0
[4,]	0	0	0	1	1	1	0	0	0
[5,]	0	0	0	1	1	1	0	0	0
[6,]	0	0	0	1	1	1	0	0	0
[7,]	0	0	0	0	0	0	1	1	1
[8,]	0	0	0	0	0	0	1	1	1
[9,]	0	0	0	0	0	0	1	1	1

or using a vector:

```
fatdiag(1:27, steps=3)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
[1,]	1	4	7	0	0	0	0	0	0
[2,]	2	5	8	0	0	0	0	0	0
[3,]	3	6	9	0	0	0	0	0	0
[4,]	0	0	0	10	13	16	0	0	0
[5,]	0	0	0	11	14	17	0	0	0
[6,]	0	0	0	12	15	18	0	0	0
[7,]	0	0	0	0	0	0	19	22	25
[8,]	0	0	0	0	0	0	20	23	26
[9,]	0	0	0	0	0	0	21	24	27

We can extract a fat diagonal and diagonalise it again.

```
m <- matrix(801:881, nrow=9, ncol=9)
fatdiag(fatdiag(m, steps=3), steps=3)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
[1,]	801	810	819	0	0	0	0	0	0
[2,]	802	811	820	0	0	0	0	0	0
[3,]	803	812	821	0	0	0	0	0	0
[4,]	0	0	0	831	840	849	0	0	0
[5,]	0	0	0	832	841	850	0	0	0
[6,]	0	0	0	833	842	851	0	0	0
[7,]	0	0	0	0	0	0	861	870	879
[8,]	0	0	0	0	0	0	862	871	880
[9,]	0	0	0	0	0	0	863	872	881

4. Conclusion

Higher-order arrays can sometimes be mapped to a matrix, which enables us to visualise these arrays in a intuitive manner. However, the standard matrix manipulations relating to diagonals become more complex when we do so. The `diagonals` package provides the `fatdiag()` function family, which enables the manipulation of fat diagonals in R, using a syntax that is very close to the `diag()` function family from R's `base` package.

Affiliation:

Bastiaan Quast
The Graduate Institute, Geneva
Maison de la paix Geneva, Switzerland
E-mail: bquast@gmail.com
URL: <http://qua.st/>